

Notes - Algorithms & Data Structures - Basics

- Dr Nick Hayward

A brief introduction to algorithms, and initial basic concepts.

Contents

- Intro
- Basic algorithms - binary search
 - conceptual example - search for a number
 - benefits of scale
- Logarithms
- Working example - binary search
 - sample - Python binary search
- Running time for algorithms
- Big O notation
 - a practical example
 - Big O usage
 - visualising big O
 - common big O runtimes
- Traveling Salesman problem
- Summary

Intro

We may consider an algorithm as a series of instructions for completing a defined task.

All code that accepts an input, and provides a defined output, may be considered analogous to an algorithm.

For example, these patterns may be seen in basic functions, which accept a parameter, and commonly return a computed value.

Algorithms come in many different shapes and sizes. For example, we commonly use algorithms with search, graphs, AI and machine learning, gaming, and many more.

For gaming, we might create an algorithm that allows the mob objects to track and follow the player using graphs. We might use *k-nearest neighbor* to define relationships with basic machine learning.

As we review and develop example algorithms, we commonly perform tests to determine performance, efficiency, speed, and comparative benefits. Such runtime testing is commonly performed using *Big-O* notation.

Basic algorithms - binary search

Binary search algorithm is a common option for finding individual items in a larger dataset.

For example, we might use this algorithm to find a person in a directory or, perhaps, find a user in a broader network.

Instead of progressing from A to B to C &c. within a defined directory, we may start in the middle and then divide the data in half.

This type of division, however, is predicated on an sort list of data for the binary search algorithm.

As binary search progresses through the dataset, it will return the index position for a matched result of `null` for no match.

This helps to eliminate possible results, and continually focus the dataset to find the search criteria.

conceptual example - search for a number

We may start with a simple example for guessing a given number from the ordered sequence 1 to 100.

For example,

- first guess is 54
 - this guess is too low
 - remove all numbers from 1 to 54
 - number sequence is updated to 55 to 100
- second guess is 75
 - this guess is too high
 - remove all numbers from 100 to 75
 - number sequence is updated to 55 to 74
- third guess is 65
 - this guess is too high
 - remove all numbers from 65 to 74
 - number sequence is updated to 55 to 64
- fourth guess is 60
 - this guess is *correct*

By using binary search, we have shown a stark contrast with the algorithm for simple search.

For example, if we consider the above number search, we can see how the algorithm optimises performance.

- 100 -> 56 -> 20 -> 10 -> 0 - answer found...

Binary search has helped us find the correct number in four turns, instead of iterating through each number sequentially.

A key part of working with binary search is the need to start with an ordered list of data.

benefits of scale

A noted benefit of this type of algorithm is the potential to scale for larger datasets.

As the dataset grows exponentially, the search algorithm is able to keep pace for simple queries.

Logarithms

We may commonly consider logs as a flipped implementation of exponentials.

For example,

- $\log_{10}100$

This represents an exponential of 2, or 10×10 . In effect, "how many 10s do we multiply together to get 100?"

For example, we may consider the following examples

exponential	logarithm
$10^2 = 100$	$\log_{10}100 = 2$
$10^3 = 1000$	$\log_{10}1000 = 3$
$2^3 = 8$	$\log_2 8 = 3$
$2^4 = 16$	$\log_2 16 = 4$
$2^5 = 32$	$\log_2 32 = 5$

Running time in Big O notation is commonly referenced as \log_2 . As we've seen, $\log_2 8 = 3$ because 2^3 gives us 8.

Likewise, for a list of 1024 elements we may test running time as $\log_2 1024$, which is the same as 2^{10} . So, a search of 1024 elements will require a maximum of 10 queries.

working example - binary search

The conceptual design and use of a binary search algorithm may be implemented in many different programming languages.

For example, we might consider the following sample for a Python application.

sample - Python binary search

The `binary_search` function takes a sorted array of items, and a single item. If the item is in the defined array, the search function will commonly return its position.

In effect, we can keep a record of where to find a given value.

So, we'll start by defining how to track the high and low values in a given data set. For example,

```
low = 0
high = len(list) - 1
```

As the example searches for a value, we keep a record of where to search in the passed array for a given value.

We may also query the middle of the array. For example,

```
mid = (low + high) / 2
guess = list[mid]
```

We may then modify these values as we use binary search with the passed dataset. For example, if we guess a value for an item, it may be higher, lower, or a known value.

For a lower value, we simply check the current stored value of `low`. If the guess is too low, we may update the current low value accordingly.

```
if (guess < item) {
    low = mid + 1
}
```

Running time for algorithms

The first option for timing algorithms is simple search. In effect, 100 items has a potential maximum number of guesses of 100.

If we increase this number exponentially, the potential maximum time will continue to grow at the same rate. So, 4 billion items may take 4 billion guesses to reach the end of the list. This is known as *linear time*.

However, if we compare this performance with binary search we quickly see the performance benefits.

For example, for a list of 100 items we require at most 7 guesses. Larger datasets see a marked improvement in performance. 4 billion results will now require a maximum of 32 guesses.

So, we have a comparative result

- $O(n)$ for linear time
- $O(\log n)$ for logarithmic time

Big O notation

Big O is special notation we may use to test and define the comparative performance of an algorithm.

For example, we commonly use this notation to test the performance of a third party algorithm. We may then compare and contrast various algorithms relative to project requirements.

a practical example

An example of choosing between simple and binary search.

n.b. this may seem like an obvious choice, but there may be contexts where *linear time* may be acceptable.

In many examples, we need an algorithm that is both fast and correct.

e.g. Landing on Mars...

We need to quickly choose an algorithm, usually in 10 seconds or less, to allow a spaceship to land on Mars.

For this test, binary search will be quicker for most tests. However, simple search is easier to write, and may reduce errors due to its inherent simplicity.

As we're performing mission critical tasks, we can't have any bugs.

So, we begin by running each algorithm 100 times. Each task may take 1 millisecond to execute. If we run initial tests, we get the following results

- simple search = 100 ms ($100 \times 1\text{ms}$)
- binary search = 7 ms ($\log_2 100 = 7$)

100 ms vs 7 ms.

Whilst the real-world usage difference is minimal, the actual program will likely require a billion plus tasks and executions.

We may perform a quick initial scaling of timings, e.g.

- binary search = ~30 ms ($\log_2 1,000,000,000$)

Binary search was initially ~15 times faster, so simple search will scale to 30×15 .

This seems reasonable, and is within tolerances for the program.

However, there's a major issue with this cursory calculation. It's based on an assumption that both search algorithms grow at the same rate.

Run times grow at different rates, thereby impacting performance relative to each dataset.

If we consider this specific example, Big O notation shows us that binary search is closer to 33 million times faster than simple search.

So, we cannot use simple search for our Mars lander.

Big O usage

Big O notation tells us the relative operations for each algorithm and dataset. In effect, how fast the algorithm is per task.

Big O notation defines binary search, for example, as $O(\log(n))$.

visualising big O

We may start with various algorithms to draw a grid of 16 squares.

We're effectively trying to determine the best algorithm to use. This decision is often predicated on many disparate conditions, which may reflect priorities in a given project.

For example, we may consider the following algorithm

- draw each box until we have the required 16 boxes...
 - big O notation tells us that a linear pattern will produce 16 grid squares
 - one box is one operation...

Big O notation produces the following results,

$O(n)$

However, there are better and more efficient options for algorithms.

For example, a second algorithm option uses folding to optimise the create of squares in the grid.

If we start by folding the a large square, a piece of paper for example, we immediately create two boxes. In effect, each fold is an operation in the algorithm.

We may continue folding the square, piece of paper, until we create our grid of 16 squares. This oly takes four operations to complete.

This algorithm produces a performance of $O(\log n)$ time.

common Big O runtimes

As we use various algorithms for projects, we may consider common performance times as calculated using Big O notation.

For example.

big O notation	description	algorithm
$O(n)$	also known as linear time	simple search
$O(\log n)$	also known as log time	binary search
$O(n * \log n)$	a faster sorting algorithm	e.g. quicksort
$O(n^2)$	a slow sorting algorithm	e.g. selection sort
$O(n!)$	a very slow algorithm	e.g traveling salesman problem

So, if we applied each algorithm to the above creation of a grid of 16 squares, we could choose the appropriate algorithm to solve the defined problem.

n.b. this is a tad simplified representation of Big O notation to the number of required operations.

We may consider the runtime and performance of an algorithm as follows,

- algorithm speed is measured using the growth in the number of required operations, and not time in seconds
- we consider the speed of increase in the runtime for an algorithm as the size of the input increase
- runtime for algorithms is defined using Big O notation
- $O(\log n)$ is faster than $O n$, and continues to get faster as the list of search items continues to increase

Traveling Salesman problem

An algorithm with a renowned bad running time.

This has become a famous problem in Computer Science, and many believe it will be very difficult to improve its performance.

The problem is as follows,

- the salesman has to visit 5 cities
- salesman wants to visit all cities in the minimum distance possible

Commonly, we might simply review each possible route and order to and from the cities.

For 5 cities, there are 120 permutations. Then, this will scale as possible

cities	permutations
6	720
7	5040
8	40320
15	1307674368000
30	265252859812191058636308480000000

For n items, it will take $n!$ (n factorial) operations to compute the result.

This is what is known as *factorial time* or $O(n!)$ time. So, as soon as the number of cities passes 100, we do not have enough time to calculate the number of permutations. The sun is forecast to collapse sooner!

Summary

We may summarise the above as follows,

- simple search is, of course, a poor choice compared to binary search
 - binary search is faster
- $O(\log n)$ is faster than $O(n)$
 - faster as search list of items grows
- do not use seconds to measure speed of algorithm
- algorithm time is measured in terms of *growth*
- *Big O* notation is used to define and record algorithm times